

CS285 - HW3

Ho Nam Nguyen

October 18, 2022

1 Basic Q-learning performance (DQN)




Figure 1

2 Double Q-learning (DDQN)




Figure 2

3 Hyperparameters - DQN




Figure 3: I vary the buffer size and observe that with too small of a buffer, transitions are getting thrown out so fast that little to no learning is possible. The default size is 50000 so decreasing this 10 times actually has little effect on the performance, i.e. orange curve still achieves a score of 100.

4 Actor-Critic sanity check




Figure 4

5 Actor-Critic on difficult tasks

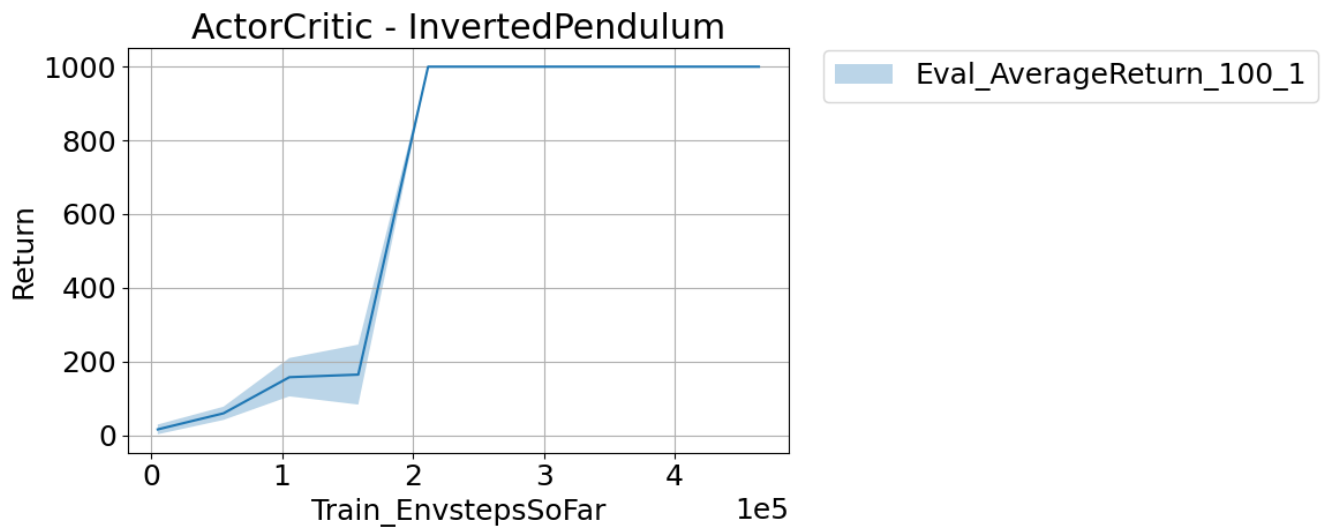


Figure 5

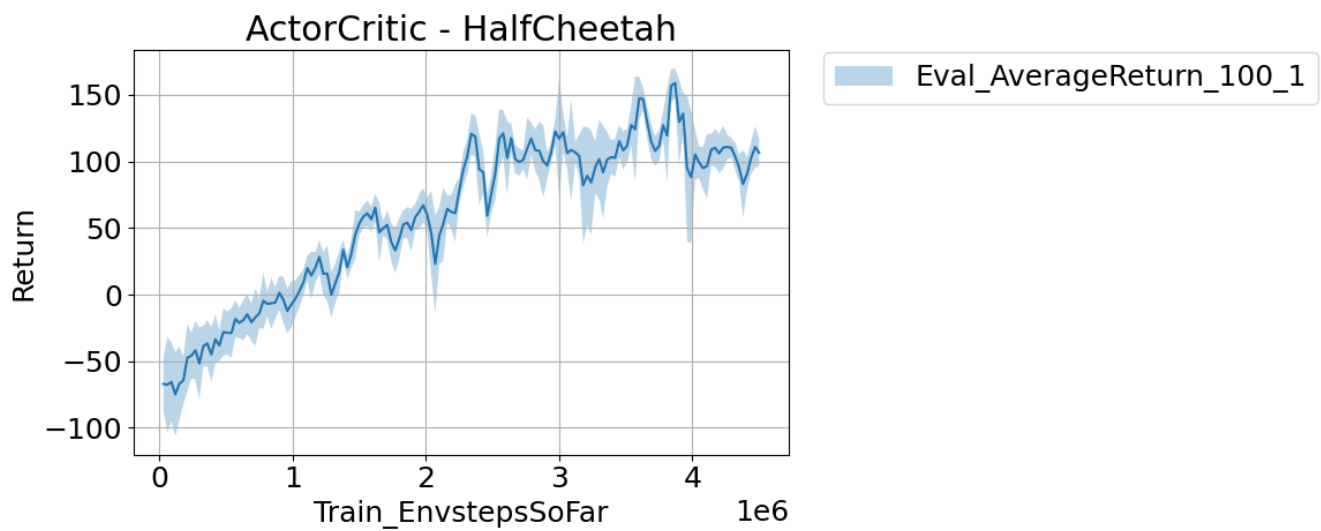


Figure 6

6 Soft Actor-Critic

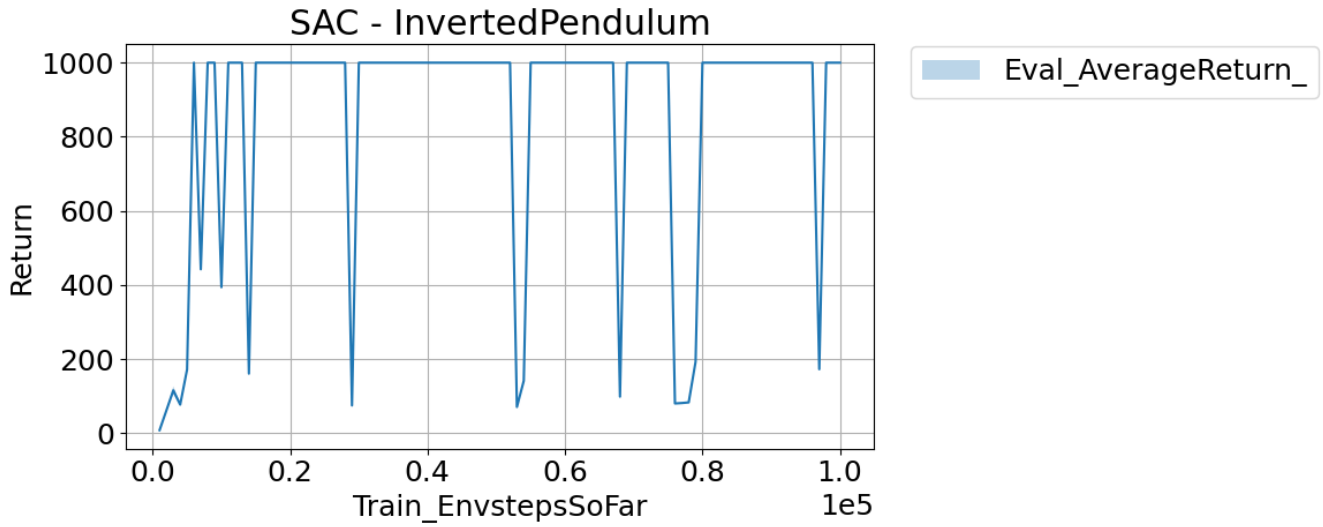


Figure 7




Figure 8